



# Apache Kafka Core Concepts

By Ramesh Fadatare (Java Guides)

# Kafka Cluster

Since Kafka is a distributed system, it act as a cluster.

A Kafka cluster consists of a set of brokers. A cluster has a minimum of 3 brokers.

Kafka Cluster



# Kafka Broker

The broker is the Kafka server. It's just a meaningful name given to the Kafka server. And this name makes sense as well because all that Kafka does is act as a message broker between producer and consumer.

The producer and consumer don't interact directly. They use Kafka server as an agent or a broker to exchange messages.



# Producer

Producer is an application that sends messages. It does not send messages directly to the recipient. It send messages only to the Kafka server.





# Consumer

Consumer is an application that reads messages from the Kafka server.

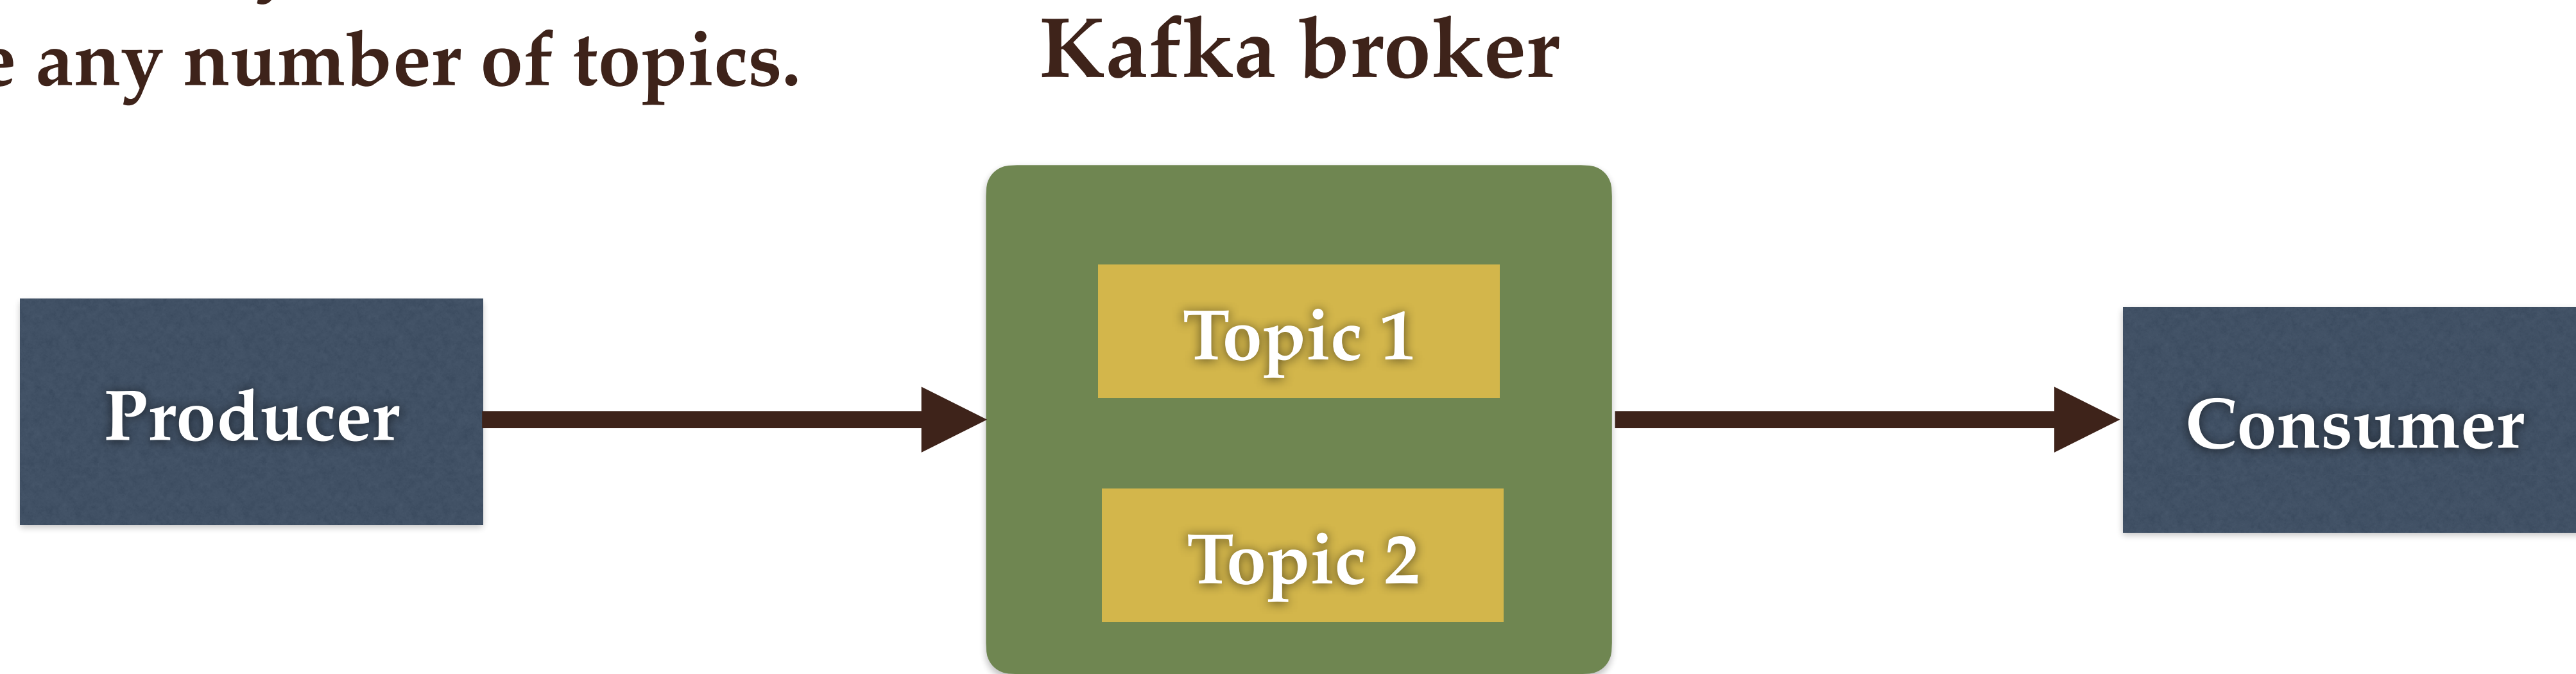
If producers are sending data, they must be sending it to someone, right? The consumers are the recipients. But remember that the producers don't send data to a recipient address. They just send it to Kafka server. And anyone who is interested in that data can come forward and take it from Kafka server. So, any application that requests data from a Kafka server is a consumer, and they can ask for data send by any producer provided they have permissions to read it.



# Kafka Topic

We learned that producer sends data to the Kafka broker. Then a consumer can ask for data from the Kafka broker. But the question is, Which data? We need to have some identification mechanism to request data from a broker. There comes the notion of the topic.

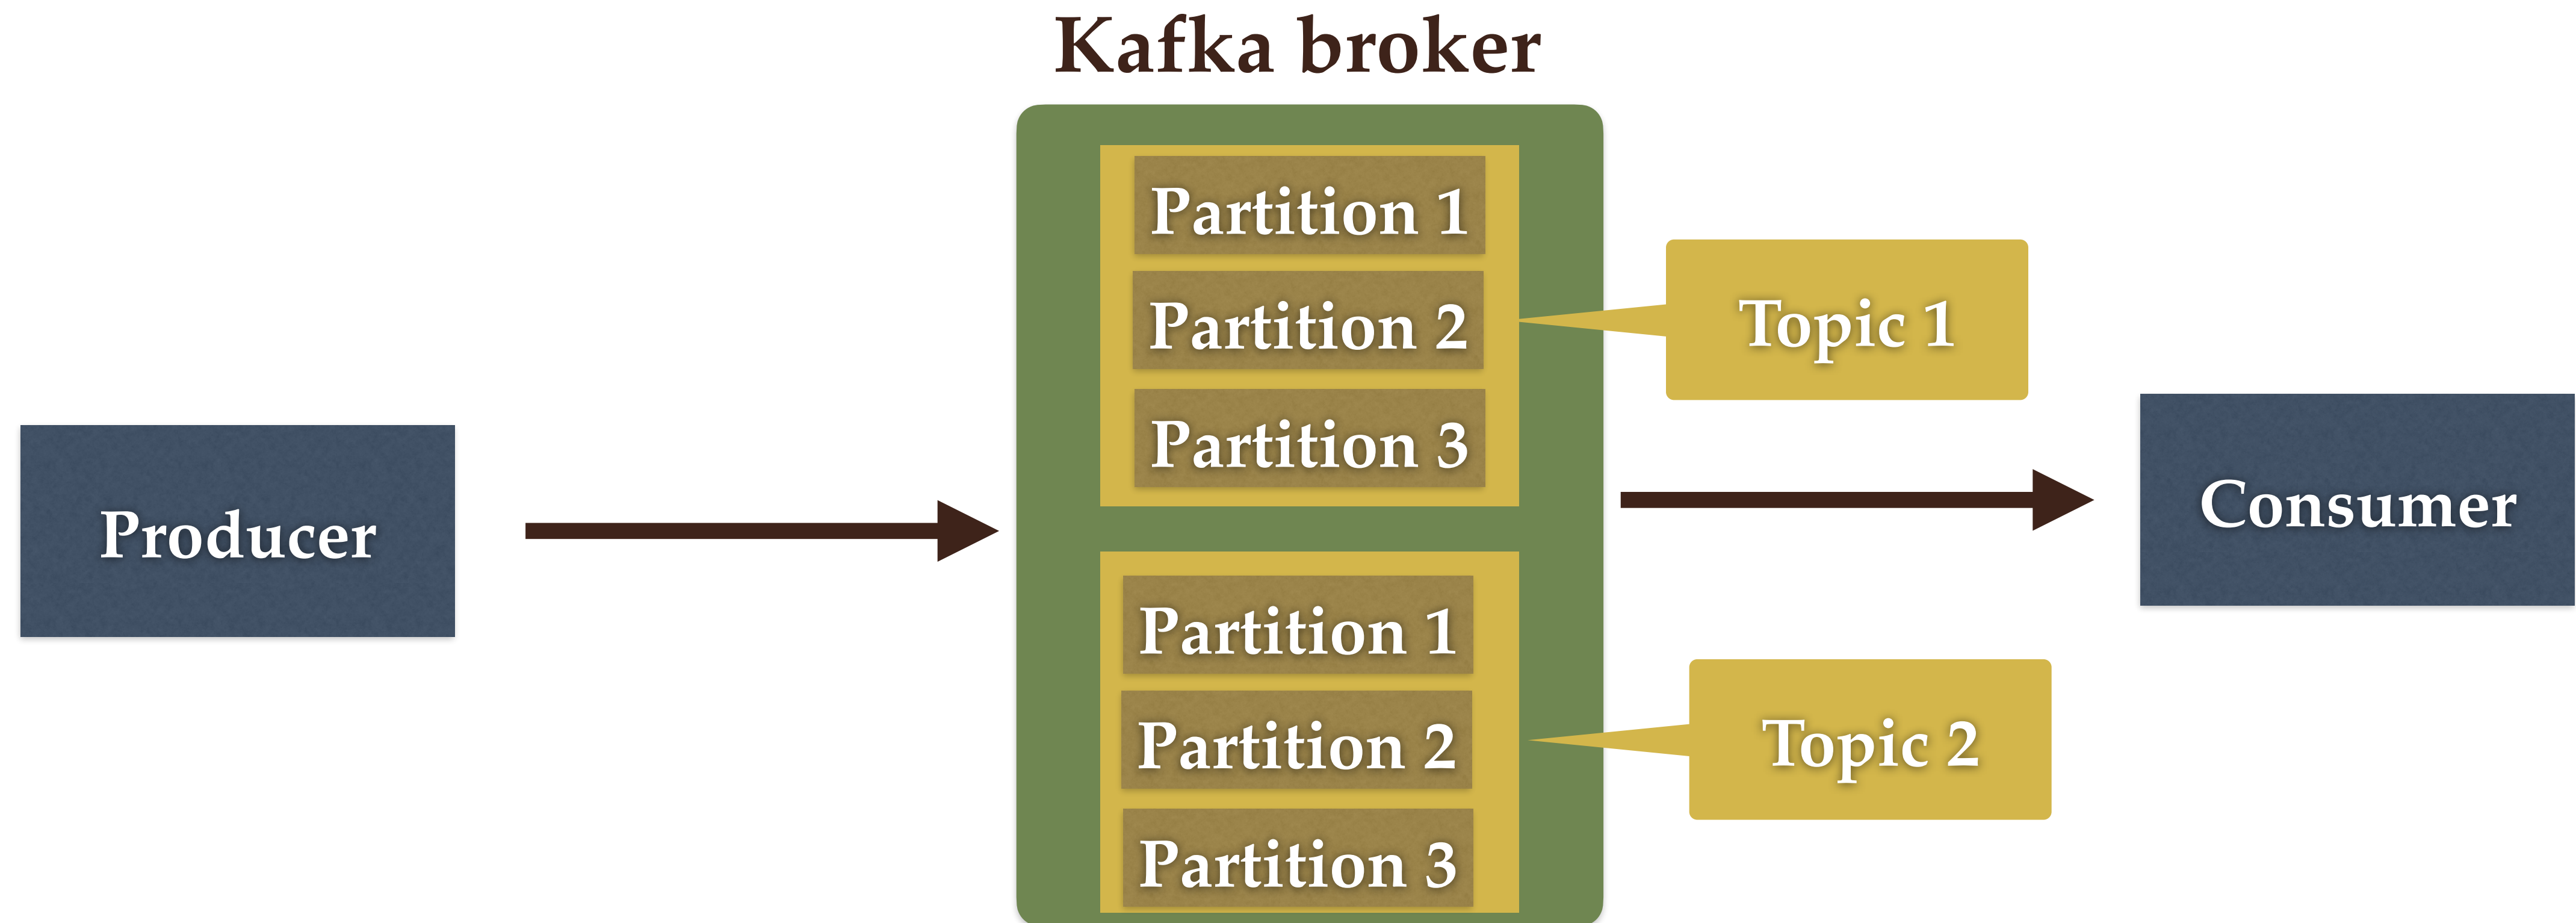
- Topic is like a table in database or folder in a file system.
- Topic is identified by a name.
- You can have any number of topics.



# Kafka Partitions

Kafka topics are divided into a number of partitions, which contain records in an unchangeable sequence.

Kafka Brokers will store messages for a topic. But the capacity of data can be enormous and it may not be possible to store in a single computer. Therefore it will be partitioned into multiple parts and distributed among multiple computers, since Kafka is a distributed system.

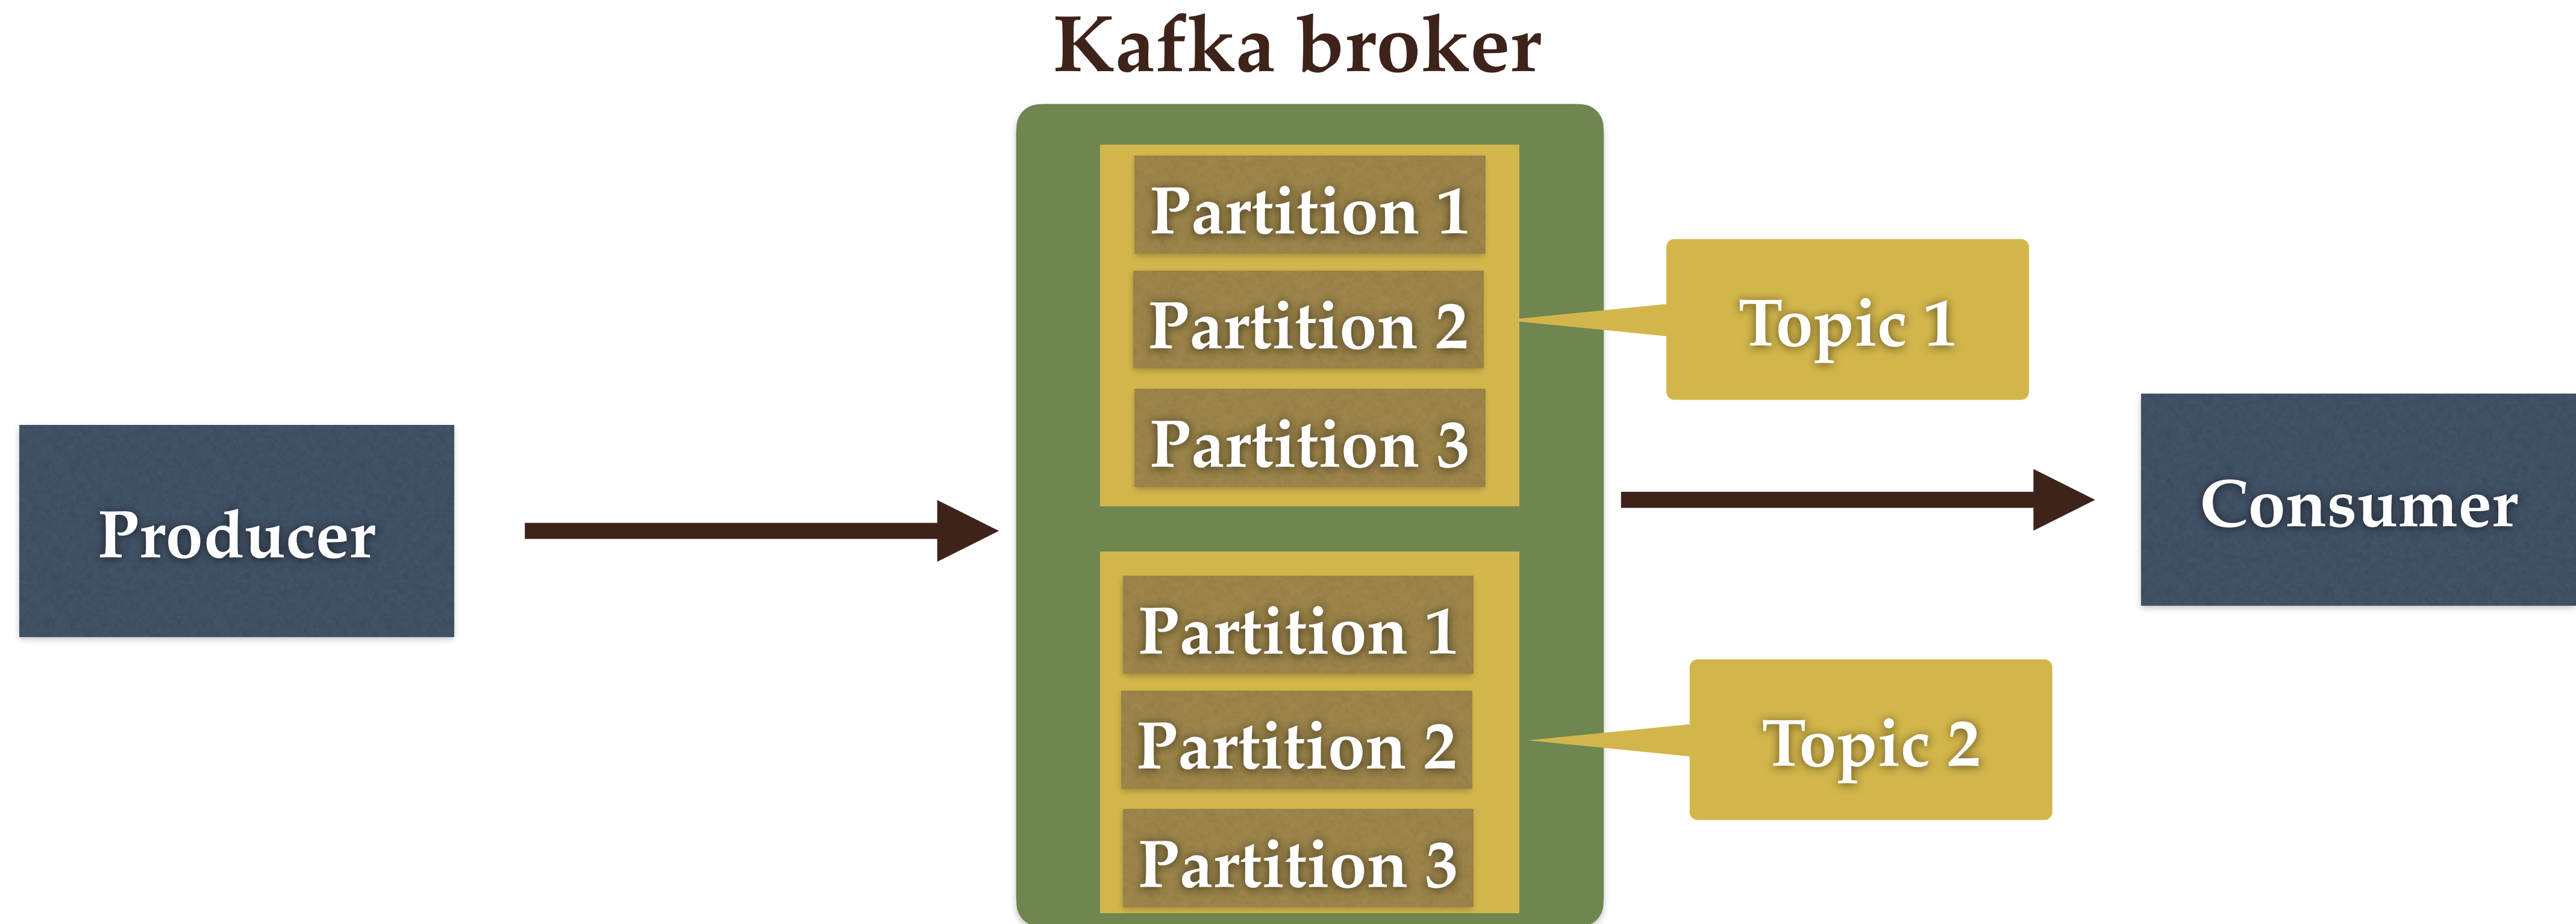




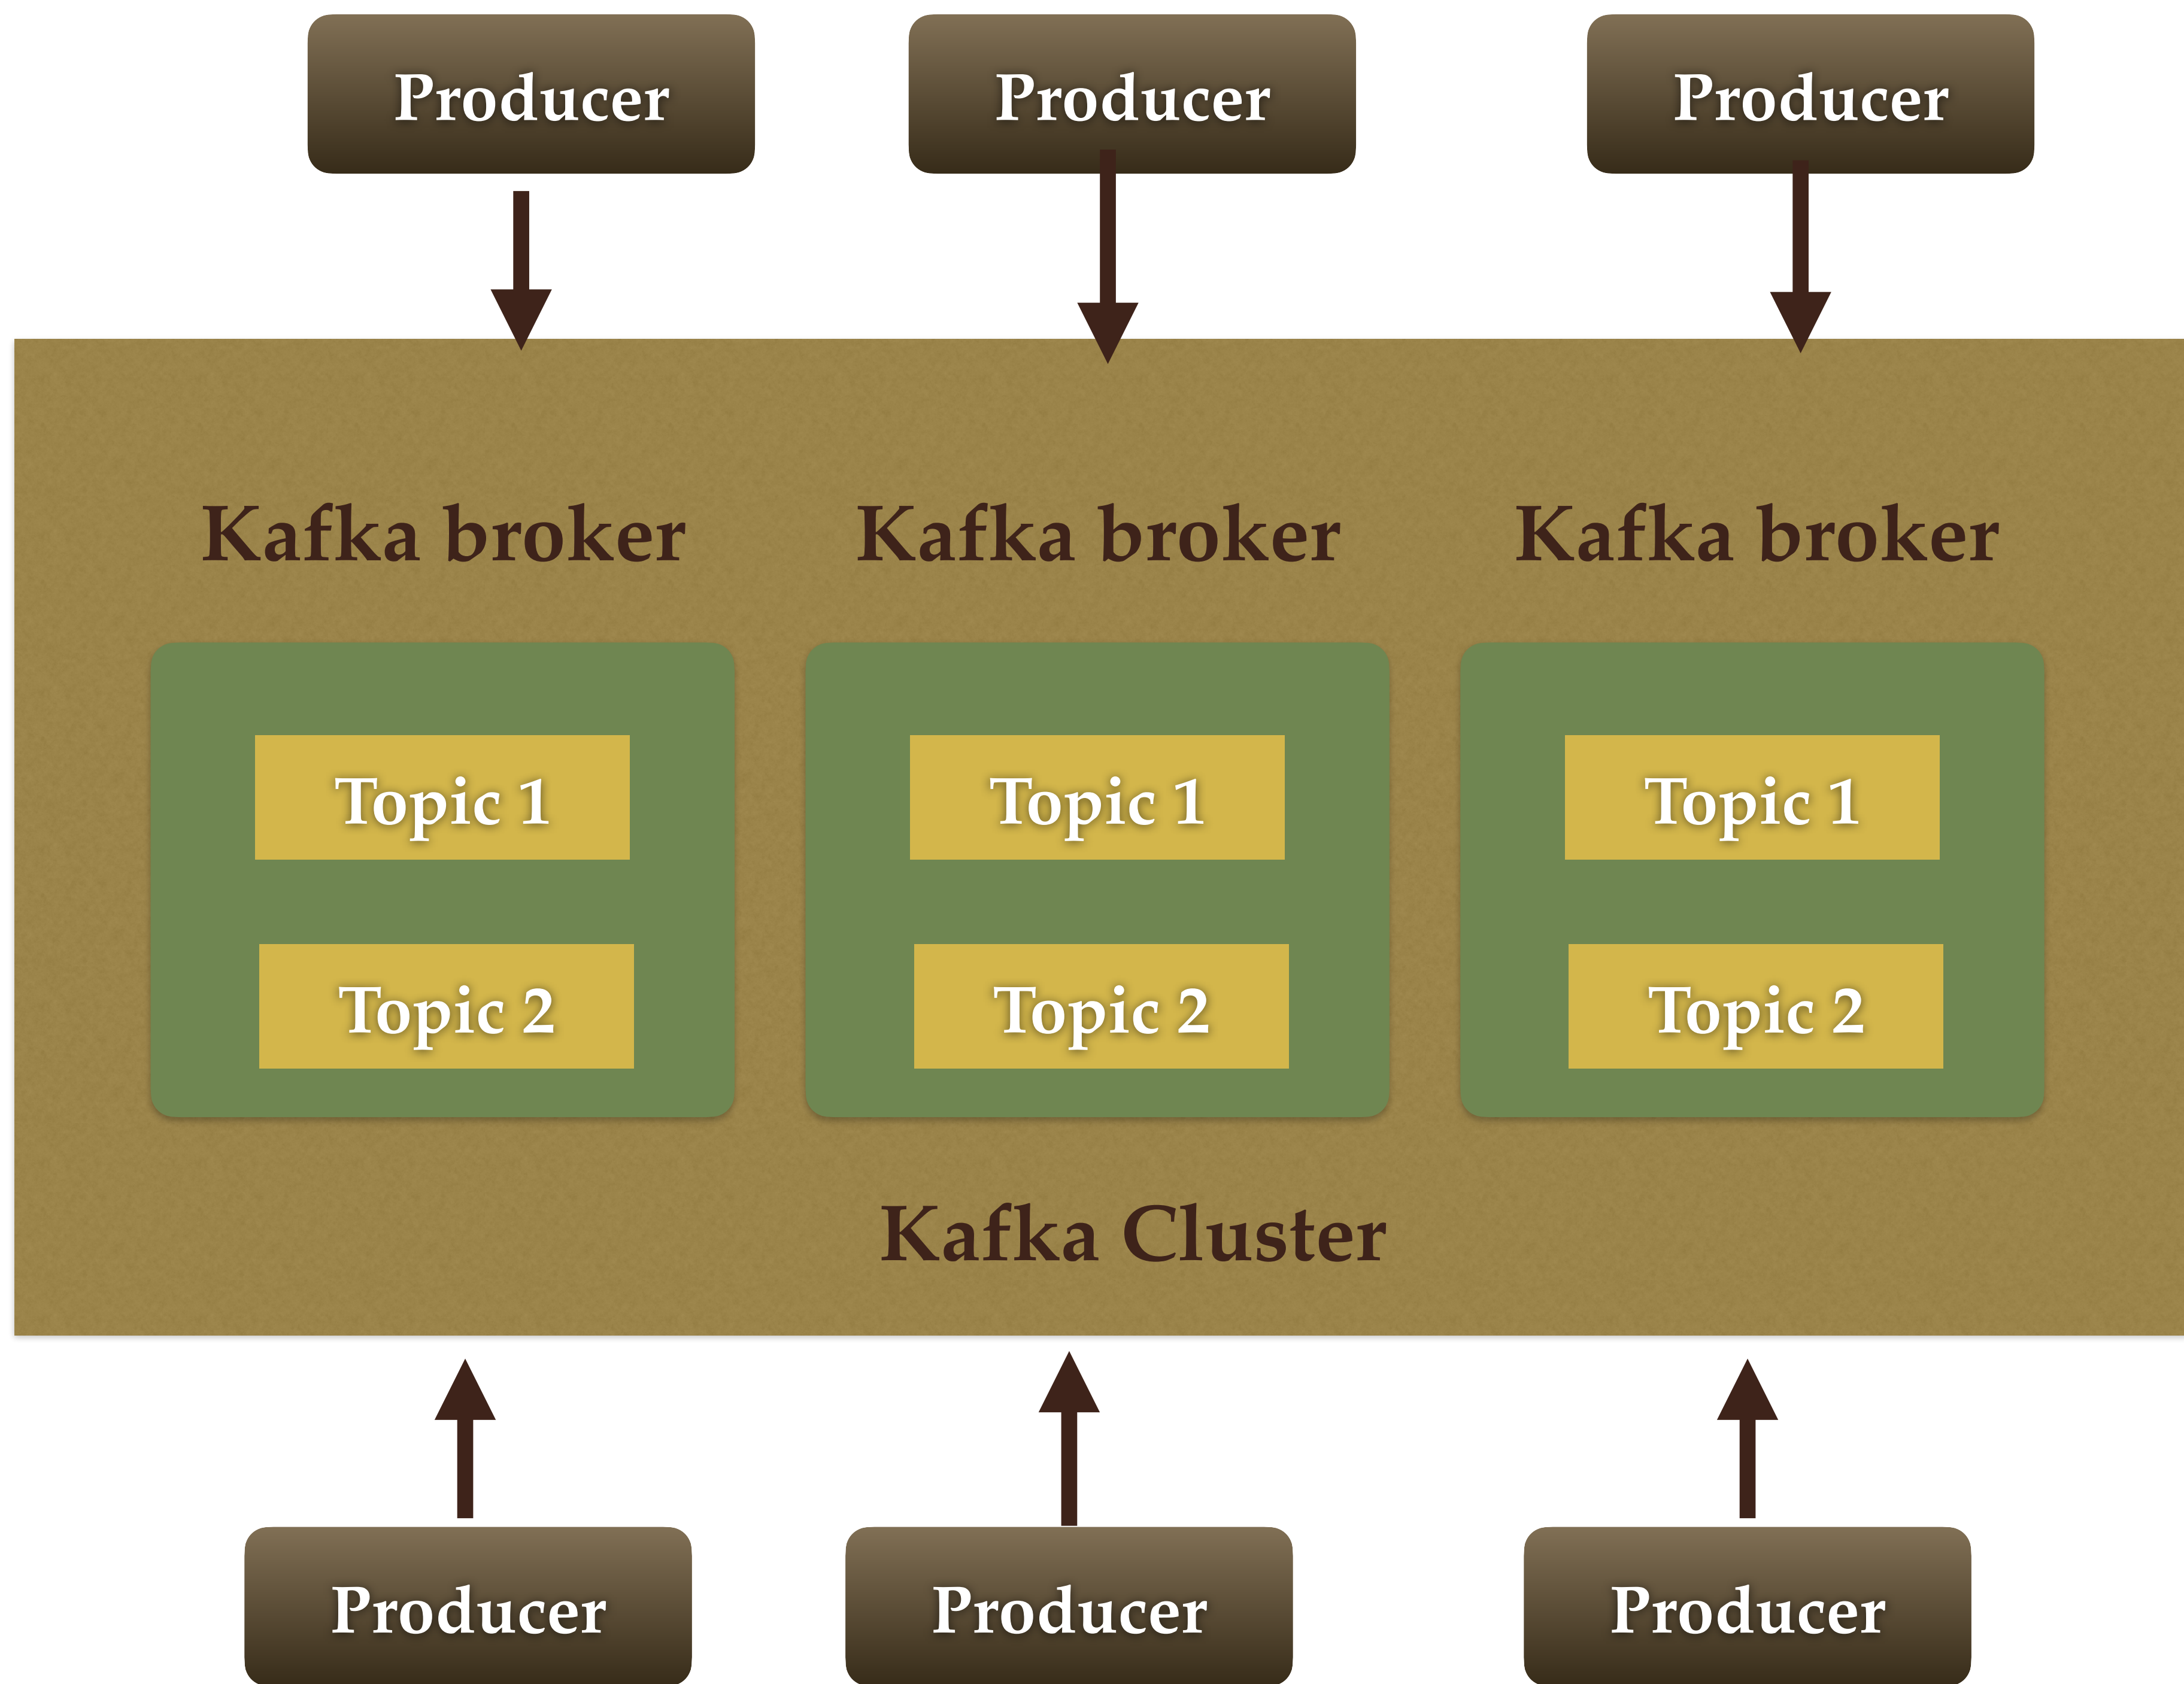
# Kafka Partitions

Kafka topics are divided into a number of partitions, which contain records in an unchangeable sequence.

Kafka Brokers will store messages for a topic. But the capacity of data can be enormous and it may not be possible to store in a single computer. Therefore it will be partitioned into multiple parts and distributed among multiple computers, since Kafka is a distributed system.







# Consumer Groups

A consumer group contains one or more consumers working together to process the messages.

